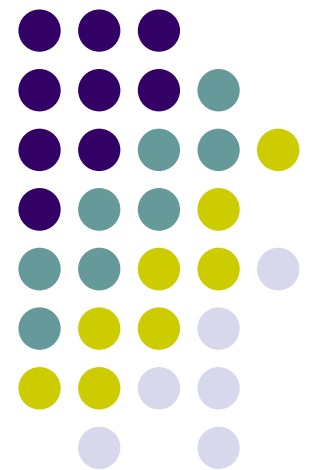


# Artificial Intelligence

CPSC 481

## Machine Learning

### Part B

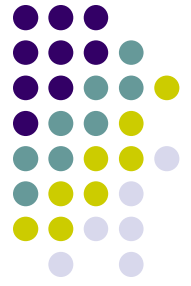




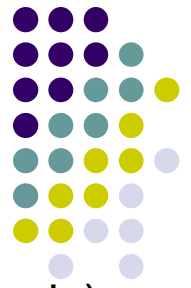
# Overview

- Non-symbolic machine learning
  - Artificial neural networks
- Unsupervised Learning
  - Clustering algorithms
- Other machine learning approaches
  - Genetic learning
  - Reinforcement learning
  - Deep learning
  - Online learning
- Issues related to machine learning

# Another Categorization of Machine Learning Approaches



- Symbolic learning vs. non-symbolic learning
  - **Symbolic learning** builds symbolic concepts or models from experience by making improvement using symbolic computation of data.
  - **Non-symbolic learning** uses optimization (using numerical computation) or adaptation (evolutionary approach)
- **Some symbolic learning approaches**
  - Concepts learning through generalization and specialization
- **Some non-symbolic learning approaches**
  - **Genetic learning** is based on **adaptation** to environment using fitness function.
  - **Artificial neural network** is based on weights optimization or gradient descent search.



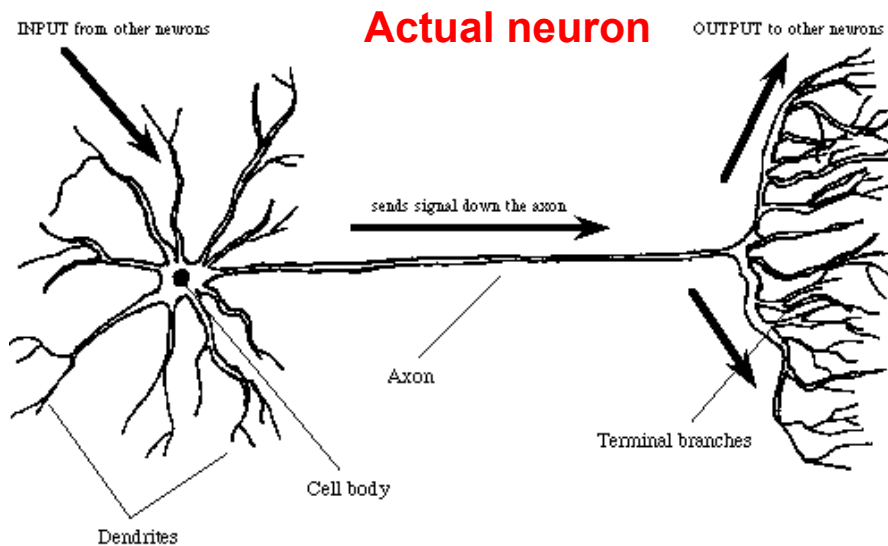
# Genetic Learning using GP

- **Learning a function using Genetic Programming** (for regression analysis)
  - Learning of the function that map a data item to a real valued prediction variable.
- **Parameters for genetic supervised learning**
  - **Objective:** evolve functions fitting the values of the cases in training data set
    - Parameters: Values
    - **Terminal set:** variable x (input) and integer constants between -5 and +5
    - **Function set:** {+, -, \*, %}
    - **Fitness function:** Error function
    - Population size: 600
    - Crossover probability: 90%
    - Mutation probability: 5%
    - Selection: tournament selection, size 4
    - Tolerant error: 0.001
    - Termination criterion: none
    - Maximum number of generations: 100
    - Maximum depth of tree after crossover: 200
    - Maximum mutant depth: 4

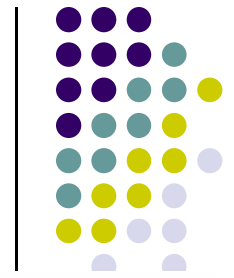
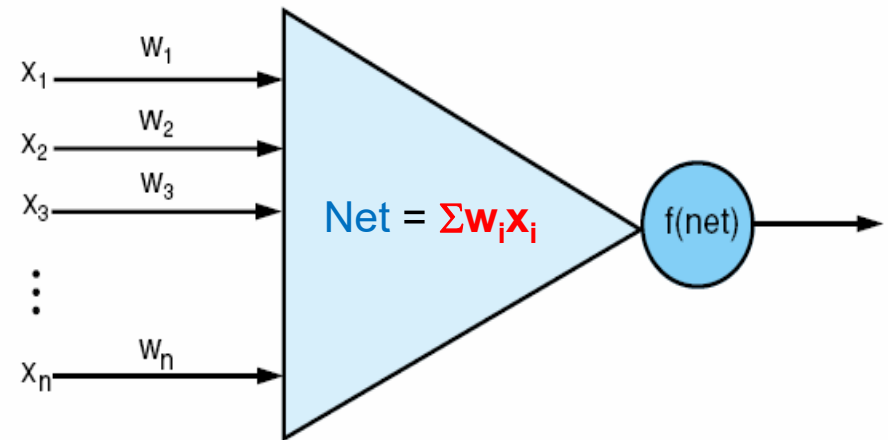
## Training data set

recordID	Input	Output
1	0.000	0.000
2	0.100	0.005
3	0.200	0.020
4	0.300	0.045
5	0.400	0.080
6	0.500	0.125
7	0.600	0.180
8	0.700	0.245
9	0.800	0.320
10	0.900	0.405

# Artificial Neuron



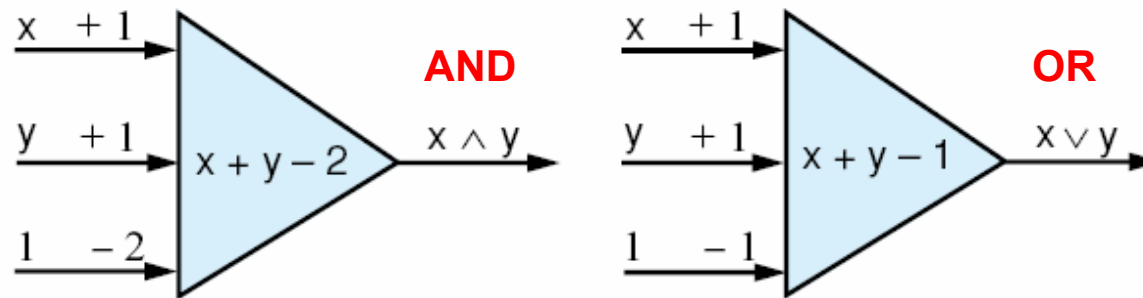
Artificial neuron



- **Artificial neuron** consists of
  - **Input signals** in a vector,  $\mathbf{X}$ : data from environment or other neurons
  - **Weights** in a vector,  $\mathbf{W}$ : connection strengths
  - **Activation function**,  $f$  that determines the neuron's output value based **net** calculated by  $\sum w_i x_i$  and a thresh hold value,  $\theta$ .
    - **Why weighted sum?**
- **Artificial neuron** is the **basic unit** of **Artificial Neural Network (ANN)**.



# McCulloch-Pitts Neuron Model to Calculate the Logic Functions, AND and OR, 1943



Bias (for constant)

Activation function  $f(\text{net})$ :

1 if  $\text{net} \geq \theta$  ( $\theta = 0$ )

0 otherwise

where  $\text{net} = \sum w_i x_i$

Input		net	AND
x	y	$x + y - 2$	Output
1	1	0	1
1	0	-1	0
0	1	-1	0
0	0	-2	0

net	OR
$X+Y-1$	Output
1	1
0	1
0	1
-1	0

What is the key idea used to solve these problems? Changing weights or activation function?

# Can we use McCulloch-Pitts Neuron Model to Solve Learning Problems?



**Class**

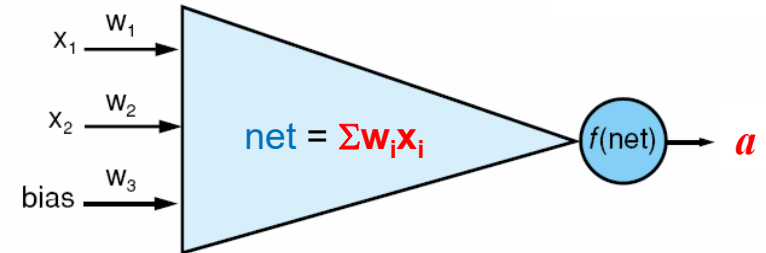
$x_1$	$x_2$	Output
1.0	1.0	1
9.4	6.4	-1
2.5	2.1	1
8.0	7.7	-1
0.5	2.2	1
7.9	8.4	-1
7.0	7.0	-1
2.8	0.8	1
1.2	3.0	1
7.8	6.1	-1

**Training data set**

- Can we use McCulloch-Pitts Neuron Model to solve other problems?
- If Yes, which part of the model do we need to change?
- How can we solve a learning (or classification) problem based on McCulloch-Pitts Neuron Model?

# Perceptron Learning

(Rosenblatt, 1962)



**Assign random values as initial weights,  $W$**

For each example **vector**,  $X^m$  from a training data set, **feed** it to the neuron as input

1. **Produce output**,  $a$  using the activation function,  $f(\text{net})$

$f(\text{net})$  returns 1 if  $\text{net} \geq \theta$  where  $\text{net} = \sum w_i x_i$   
else -1 (if  $\text{net} < \theta$ )

2. **Adjust weights** for next example,  $X^{m+1}$  as follows:

If  $t = a$  use the current weights,  $W^m$  (or do nothing)

where  $t$ : the desired output in example  $X^m$  in vector

**Else** (error) adjust weights

- Compute  $\Delta W^m = l(t - a)X^m$ , where  $l$  is a **learning rate** (0.0~1.0) to control the speed of convergence
- Weights for next example  $X^{m+1}$ ,  $W^{m+1} = W^m + \Delta W^m$

**Repeat** this process for all examples in the training set **until converged** (possibly reusing the training data).

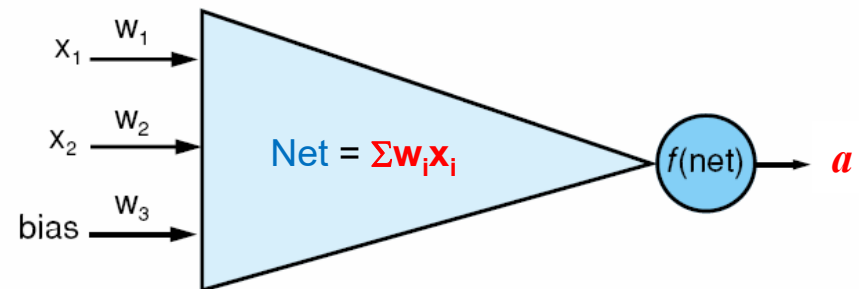


# Example: Perceptron Learning



		Class
$x_1$	$x_2$	Output
1.0	1.0	1
9.4	6.4	-1
2.5	2.1	1
8.0	7.7	-1
0.5	2.2	1
7.9	8.4	-1
7.0	7.0	-1
2.8	0.8	1
1.2	3.0	1
7.8	6.1	-1

Training data set



$$f(\text{net}) = f(w_1 * x_1 + w_2 * x_2 + w_3 * 1), \text{ bias}=1, \theta = -0.6$$

$$\mathbf{W} = [0.75, 0.5, -0.6] \text{ (randomly generated), } l = 0.2$$

$$f(\text{net})^1 = f(0.75 * 1 + 0.5 * 1 - 0.6 * 1) = f(0.65) = 1,$$

$$(t - a) = (1 - 1) = 0, \text{ so no } \Delta w \text{ needed}$$

$$f(\text{net})^2 = f(0.75 * 9.4 + 0.5 * 6.4 - 0.6 * 1) = f(9.65) = 1,$$

$$(t - a) = (-1 - 1) = -2, \text{ we need } \Delta w$$

$$\mathbf{W}^3 = \mathbf{W}^2 + 0.2(-1 - 1)\mathbf{X}^2 = \begin{bmatrix} 0.75 \\ 0.50 \\ -0.60 \end{bmatrix} - 0.4 \begin{bmatrix} 9.4 \\ 6.4 \\ 1.0 \end{bmatrix} = \begin{bmatrix} -3.01 \\ -2.06 \\ -1.00 \end{bmatrix}$$

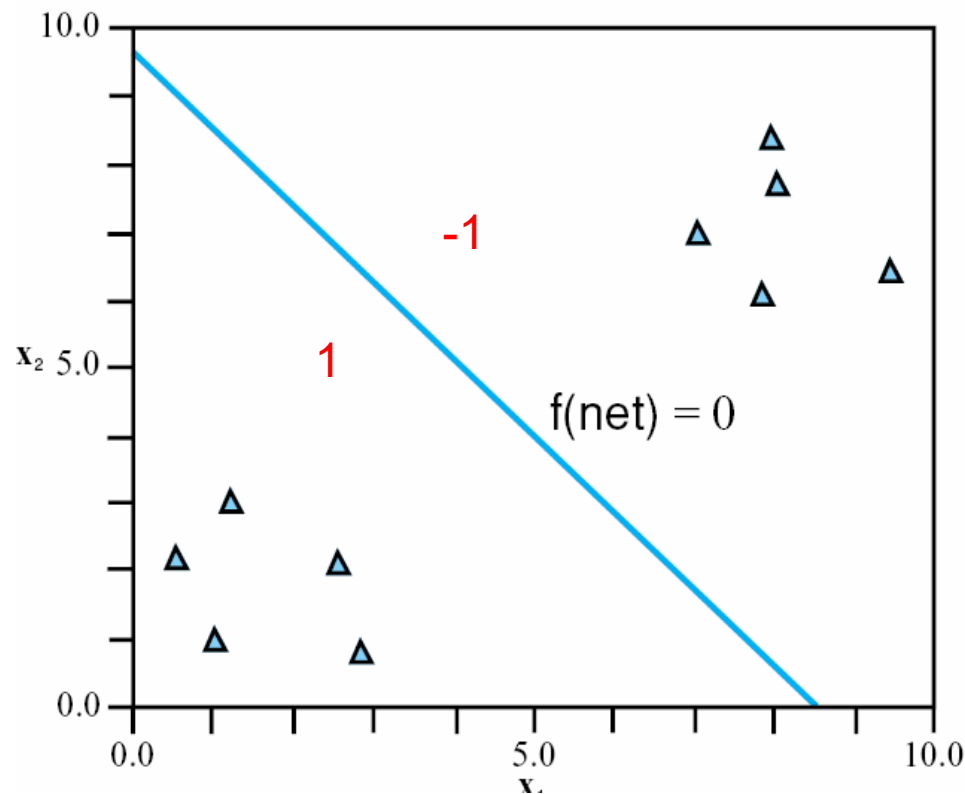
$$f(\text{net})^3 = f(-3.01 * 2.5 - 2.06 * 2.1 - 1.0 * 1) = f(-12.84) = -1, \text{ we need } \Delta w$$

$$\mathbf{W}^4 = \mathbf{W}^3 + 0.2(1 - (-1))\mathbf{X}^3 = \begin{bmatrix} -3.01 \\ -2.06 \\ -1.00 \end{bmatrix} + 0.4 \begin{bmatrix} 2.5 \\ 2.1 \\ 1.0 \end{bmatrix} = \begin{bmatrix} -2.01 \\ -1.22 \\ -0.60 \end{bmatrix}$$

\*Continue this process until converged.

(This is an example of gradient descent search.)

# The Perceptron that Linearly Separates the Training Data Set



A linearly separable function

\*The weight vector converged after repeated training on the data set, about 500 iterations

$$W^{500} = [-1.3, -1.1, 10.9]$$

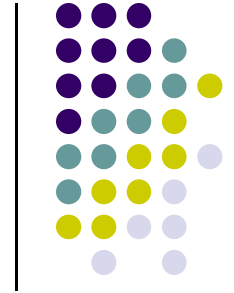
$f(\text{net}) = w_1x_1 + w_2x_2 + w_3$ , where two discriminant functions  $g_i$  (1) and  $g_j$  (-1) are the same (border),  $g_i = g_j$

$$f(\text{net}) = -1.3x_1 - 1.1x_2 + 10.9 = 0$$

$f(\text{net})$  is the line separating the two classes that can be used for classification.

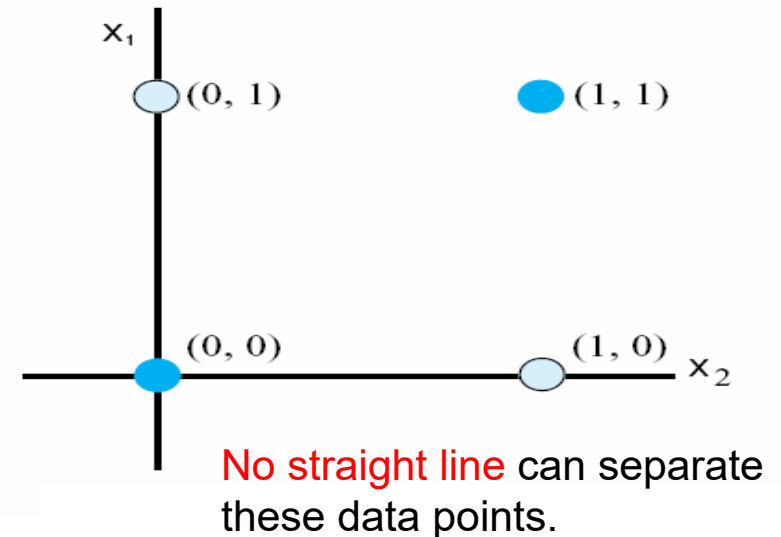
+How is the perceptron learning different from symbolic learning approaches such as Winston's learning program and decision tree induction?<sub>10</sub>

# Linearly Not Separable Case



Truth table for **exclusive-OR**

$x_1$	$x_2$	Output
1	1	0
1	0	1
0	1	1
0	0	0



From the data set, **following equations** must be solved:

$w_1 * 1 + w_2 * 1 < \theta$  , where  $\theta$  is a threshold.

$w_1 * 1 + 0 > \theta$

$0 + w_2 * 1 > \theta$

$0 + 0 < \theta$  or  $\theta$  must be positive

(This series of equations with  $w_1$ ,  $w_2$ , and  $\theta$  has no solution)

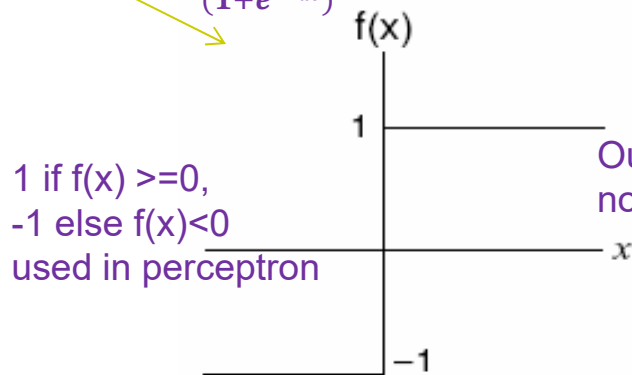
\*Minsky and Papert (1969) proved that a **perceptron** that solves **exclusive-or** is **impossible**.

# Sigmoidal Function for Activation Function

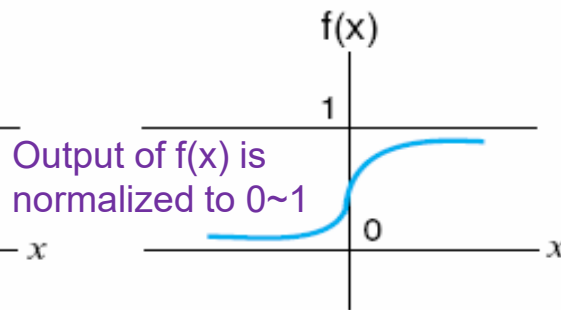


## Sigmoidal function:

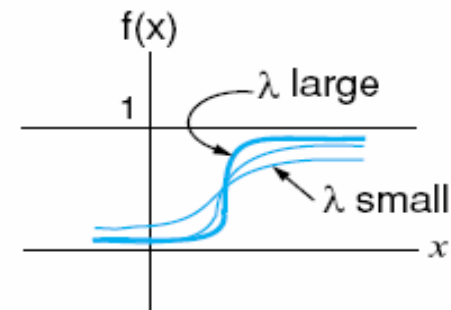
$f(x) = \frac{1}{(1+e^{-\lambda x})}$ ,  $\lambda$  is a **squashing parameter** to fine-tune the curve by adjusting the slope.



a. A hard limiting and bipolar linear threshold.



b. A sigmoidal and unipolar threshold.



c. The sigmoidal, biased and squashed. As  $\lambda$  gets larger the sigmoid approximates a linear threshold.

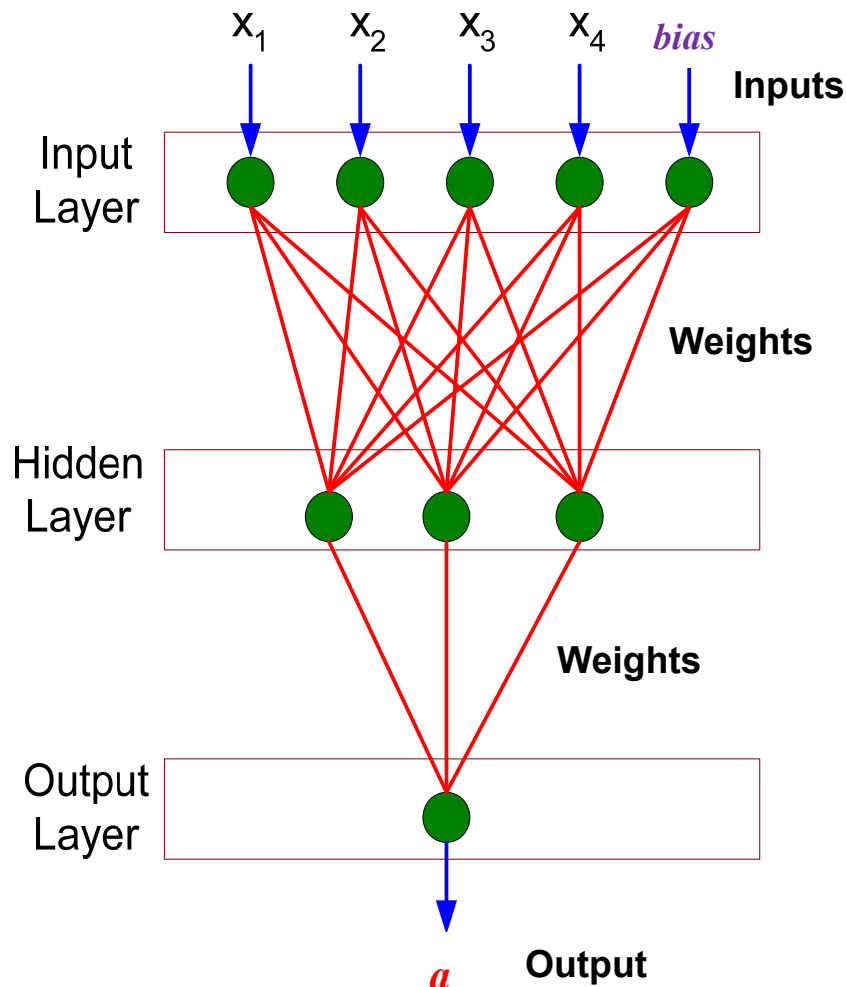
## Useful features of sigmoidal function:

- S-shaped non-linear and continuous function in the range of 0 ~1
- Differentiable (has a derivative) at each point and derivative calculation is easy.
- The value of derivative is greatest where the sigmoidal function is changing most rapidly.

**Sigmoidal activation function:**  $f(\text{net}) = 1/(1 + e^{-\lambda \text{net}})$ , where  $\text{net} = \sum w_i x_i$

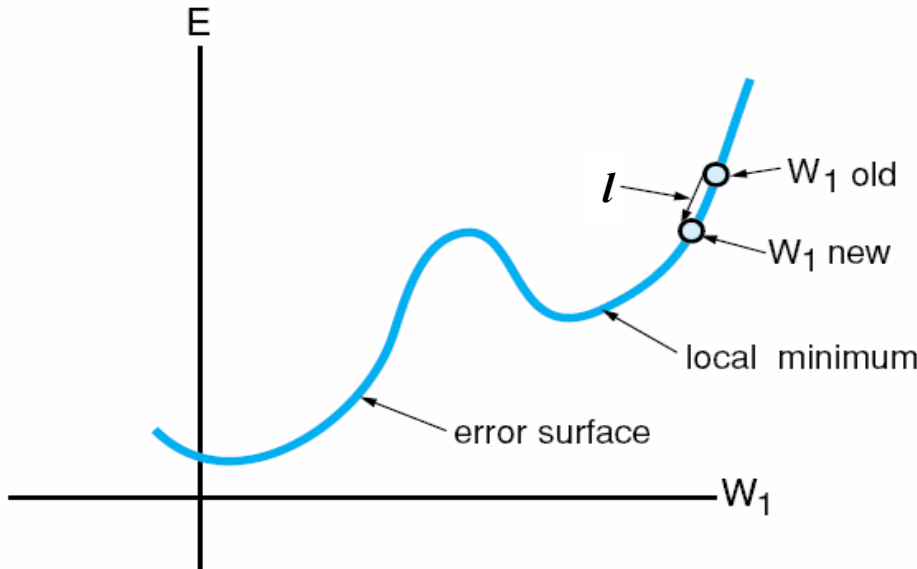
The value of  $f(\text{net})$  is continuous 0.0~1.0, not binary.

# Multi-layer Neural Networks with Backpropagation



- **Goal:**
  - Given training data, **find weights** that **fit** outputs of neural net to desired output or classification.
- **Mechanism:**
  - For each training example, **propagate the inputs Forward** for activation.
  - **Backpropagate** the error to previous layer and **Adjust** the weights slowly in direction that reduces error.
  - **Repeat** this process until converged.
- **Proven to converge**  
on the nearest **local minimum** but
  - May not find the best weights
  - Sensitive to initial weight setting
  - Long training time
- **Many applications:**  
solved many difficult real-world problems.

# Weight Adjustment with Sigmoidal Function



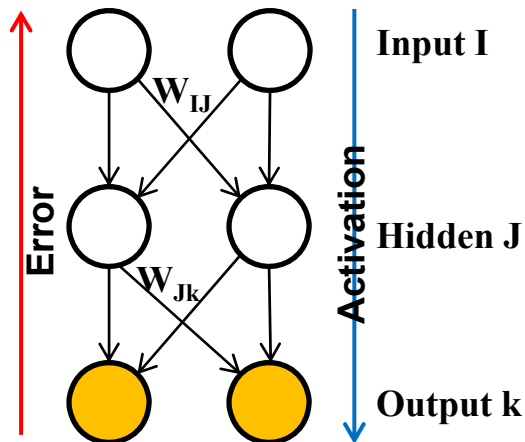
\*The **error surface** represents **cumulative error** over a data set as a function of weights.

\* $l$  determines the size of learning step.  
\***Gradient** is a measure of **slope**, as a **function of direction** from a point on a surface.  
\*Learning algorithm uses **gradient descent** method to find the **direction** on the **error surface** which most rapidly reduces the error (ending at a local minimum).

How do we determine + or - of  $\Delta W$ ?

- **Weight adjustment by the Delta rule** (Rumelhart et al. 1986)
  - **Weight adjustment** on the  $j$ th node from the  $i$ th input:  $\Delta W_j = l(t_j - a_j)f'(\text{net}_j)x_i$  where  $l$  is the learning rate;  $t_j$  is the desired output;  $a_j$  is the output of the node;  $f'(\text{net}_j)$  is the derivative of  $f(\text{net})$ ;  $x_i$  is an input (or output from the previous node)
  - Why  $f'(\text{net})$ ? How do we calculate  $f'(\text{net})$ ?
  - Derive the **delta rule**!

# Computing $\Delta w$ for Output Nodes



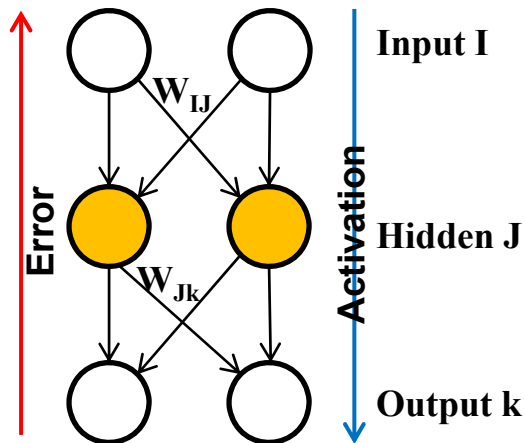
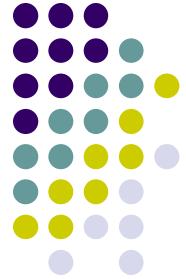
## Symbols

- **I:** input layer
- **J:** hidden layer
- **K:** output layer
- $w_{ij}$ : weight, input  $\rightarrow$  hidden
- $w_{jk}$ : weight, hidden  $\rightarrow$  output
- $a$ : activation value, output
- $t$ : target value
- $net$ : the net input
- $l$ : learning rate

## Gradient descent on error

- $E = (1/2) \sum_k (t_k - a_k)^2$  the total error in a network
- $\Delta w \propto -\frac{\partial E}{\partial w}$  to adjust the weights to **reduce** the overall error
- $\Delta w_{jk} \propto -\frac{\partial E}{\partial w_{jk}}$  to adjust the weight at the output layer
- $\Delta w_{jk} = -l \frac{\partial E}{\partial a_k} * \frac{\partial a_k}{\partial net_k} * \frac{\partial net_k}{\partial w_{jk}}$  since error is not directly a function of a weight
- Derivative of the error with respect to the activation
  - $\frac{\partial E}{\partial a_k} = \frac{\partial (1/2)(t_k - a_k)^2}{\partial a_k} = -(t_k - a_k)$
- Derivative of the activation with respect to the net input
  - $\frac{\partial a_k}{\partial net_k} = f'(net) = \partial(1 + e^{-net_k})^{-1} / \partial net_k = a_k(1 - a_k)$
- Derivative of the net input with respect to a weight
  - $\frac{\partial net_k}{\partial w_{jk}} = \frac{\partial (w_{jk} a_j)}{\partial w_{jk}} = a_j$
- Weight change rule (delta rule) for a hidden to output weight
  - $\Delta w_{jk} = -l \frac{\partial E}{\partial w_{jk}} = -l \frac{\partial E}{\partial a_k} * \frac{\partial a_k}{\partial net_k} * \frac{\partial net_k}{\partial w_{jk}} = l(t_k - a_k) a_k(1 - a_k) a_j$
  - If we let  $\delta_k = (t_k - a_k) a_k(1 - a_k)$ , then  $\Delta w_{jk} = l \delta_k a_j$ <sup>15</sup>

# Computing $\Delta w$ for Hidden Nodes



## Subscript notations

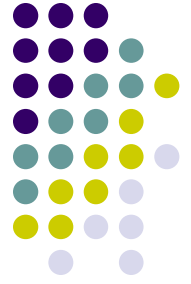
- **I**: input layer
- **J**: hidden layer
- **K**: output layer
- $w_{ij}$ : weight, input  $\rightarrow$  hidden
- $w_{jk}$ : weight, hidden  $\rightarrow$  output
- $a$ : an activation value
- $t$ : a target value
- **net**: the net input
- $l$ : learning rate

## Gradient descent on error

- $\Delta w_{ij} \propto -\frac{\partial E}{\partial w_{ij}}$  to adjust the weight at the output layer
- $\Delta w_{ij} = -l \frac{\partial E}{\partial a_j} * \frac{\partial a_j}{\partial net_j} * \frac{\partial net_j}{\partial w_{ij}}$
- Derivative of the output at J layer with respect to the input I
  - $\frac{\partial a_j}{\partial net_j} * \frac{\partial net_j}{\partial w_{ji}} = a_j(1 - a_j) * \frac{\partial net_j}{\partial w_{ji}} = a_j(1 - a_j) * \frac{\partial(w_{ji}a_i)}{\partial w_{ji}} = a_j(1 - a_j)a_i$
- Derivative of the error at K layer with respect to the activation of J layer
  - $\frac{\partial E}{\partial a_j} = [\sum \frac{\partial E}{\partial a_k} * \frac{\partial a_k}{\partial net_k} * \frac{\partial net_k}{\partial a_j}]$  since error on output is propagated from hidden layer (all nodes at J)
  - $\frac{\partial E}{\partial a_k} = -(t_k - a_k)$
  - $\frac{\partial a_k}{\partial net_k} = a_k(1 - a_k)$
  - $\frac{\partial net_k}{\partial a_j} = \frac{\partial(w_{jk}a_j)}{\partial a_j} = w_{jk}$
- Weight change rule (delta rule) for a hidden to output weight
  - Let  $\delta_k = \sum_k (t_k - a_k) a_k (1 - a_k) w_{jk}$
  - Let  $\delta_j = \delta_k a_j (1 - a_j)$
  - $\Delta w_{ij} = -l \frac{\partial E}{\partial a_j} * \frac{\partial a_j}{\partial net_j} * \frac{\partial net_j}{\partial w_{ij}} = l \delta_j a_i$

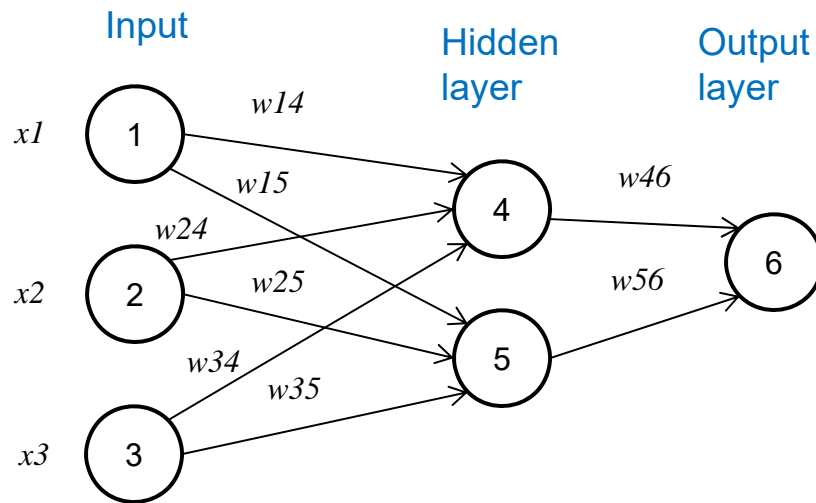
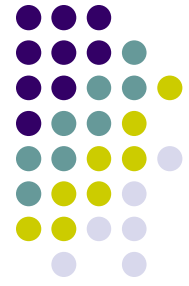


# Backpropagation Algorithm



- **Input:** Training data set; learning rate,  $l$ ; multi-layer network
- **Output:** A neural network trained to classify the data
- **Algorithm:**
  - 1) Initialize all weights and biases in network
  - 2) while terminating condition is not satisfied
  - 3)   for each sample  $X$  in training data set
  - 4)     for each hidden or output layer unit  $j$
  - 5)          $I_j = \sum_i w_{ij}a_i + \theta_j$ ; //compute the input of unit  $j$ , **net<sub>j</sub>** with respect to the previous layer  $i$ .
  - 6)          $a_j = \frac{1}{1+e^{-\lambda I_j}}$ ; //compute the output of each unit  $j$ ; assume  $\lambda = 1$ .
  - 7)     for each unit  $j$  in the output layer
  - 8)          $E_j = a_j(1 - a_j)(t_j - a_j)$ ; //compute the error to backpropagate the errors
  - 9)     for each unit  $j$  in the hidden layer
  - 10)          $E_j = a_j(1 - a_j)\sum_k E_k w_{jk}$ ; //compute the error with respect to the next layer,  $k$
  - 11)     for each weight  $w_{ij}$  in network
  - 12)          $\Delta w_{ij} = l * E_j * a_i$ ; //weight increment;  $l$  is the learning rate.
  - 13)          $w_{ij} = w_{ij} + \Delta w_{ij}$ ; //weight update
  - 14)     for each bias  $\theta_j$  in network
  - 15)          $\Delta \theta_j = l E_j$ ; //bias increment
  - 16)          $\theta_j = \theta_j + \Delta \theta_j$ ; //bias update

# Example: Backpropagation Calculation



Initial input and class

$x1$	$x2$	$x3$	class
1	0	1	1

biases

$\theta_4$	$\theta_5$	$\theta_6$
-0.4	0.2	0.1

Initial weights,  $\lambda=0.9$

$w14$	$w15$	$w24$	$w25$	$w34$	$w35$	$w46$	$w56$
0.2	-0.3	0.4	0.1	-0.5	0.2	-0.3	-0.2

Inputs and outputs for node 1, 2, 3 or  $I_1, I_2, I_3$  are the same as initial inputs  $x1, x2$ , and  $x3$ , respectively. The input and output calculations for nodes 4, 5, 6 are given below:

Unit $j$	Input, $I_j = \sum_i w_{ij}a_i + \theta_j$	Output, $a_j = \frac{1}{1+e^{-\lambda I_j}}, \lambda=1$
4	$I_4 = 0.2 + 0 - 0.5 - 0.4 = -0.7$	$1/(1+e^{0.7}) = 0.332$
5	$I_5 = -0.3 + 0 + 0.2 + 0.2 = 0.1$	$1/(1+e^{-0.1}) = 0.525$
6	$I_6 = (-0.3)(0.332) - (0.2)(0.525) + 0.1 = -0.105$	$1/(1+e^{0.105}) = 0.474$

Unit j	$E_j = a_j(1 - a_j)(t_j - a_j)$ for output layer $E_j = a_j(1 - a_j)\sum_k E_k w_{jk}$ for hidden layer
6 (output)	$(0.474)(1-0.474)(1-0.474) = 0.1311$
5 (hidden)	$(0.525)(1-0.525)(0.1311)(-0.2) = -0.0065$
4 (hidden)	$(0.332)(1-0.332)(0.1311)(-0.3) = -0.0087$

Calculation of the error at each node



Weight/bais	Updated weights by $\Delta w_{ij} = l E_j a_i$ ; $w_{ij} = w_{ij} + \Delta w_{ij}$ ; $\Delta \theta_j = l E_j$ ; $\theta_j = \theta_j + \Delta \theta_j$ ; $l=0.9$
w46	$-0.3+(0.9)(0.1311)(0.332) = -0.261$
w56	$-0.2+(0.9)(0.1311)(0.525) = -0.138$
w14	$0.2+(0.9)(-0.0087)(1) = 0.192$
w15	$-0.3+(0.9)(-0.0065)(1) = -0.306$
w24	$0.4+(0.9)(-0.0087)(0) = 0.4$
w25	$0.1+(0.9)(-0.0065)(0) = 0.1$
w34	$-0.5+(0.9)(-0.0087)(1) = -0.508$
w35	$0.2+(0.9)(-0.0065)(1) = 0.194$
$\theta_6$	$0.1+(0.9)(0.1311) = 0.218$
$\theta_5$	$0.2+(0.9)(-0.0065) = 0.194$
$\theta_4$	$-0.4+(0.9)(-0.0087) = -0.408$

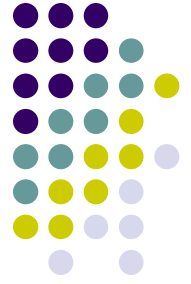
Calculations for weight and bias update

**Exercise:** Show the process of solving Exclusive-OR problem.



# Various Types of Neural Networks

- **Different neural networks can be built using different options:**
  - Network topology
    - Many layers instead of 2 => +100 layers (**deep learning**)
  - Learning algorithm
  - Encoding scheme
- **Neural network can be used for both**
  - Supervised learning (classification)
  - Unsupervised learning (clustering)



# Unsupervised Learning

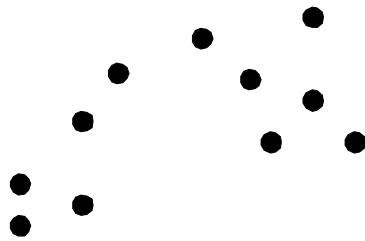
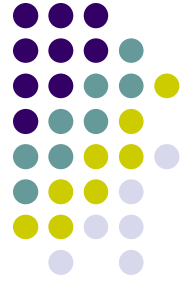
- Unsupervised learning eliminates the teacher and requires that the learners learn concepts on their own
  - **AM** was the first unsupervised discovery program to learn mathematical concepts, e.g., concepts of set theory, natural number theory, prime numbers, etc. but was not very successful, mostly due to representation (symbolic based) and inability to learn.
- **Methods needed for unsupervised learning**
  - **Discovery** such as scientific discovery process
    - hypothesis → generality → conclusion
  - **Pattern finding**
    - **Pattern recognition**
    - **Clustering** (class finding by similarity)

# Clustering

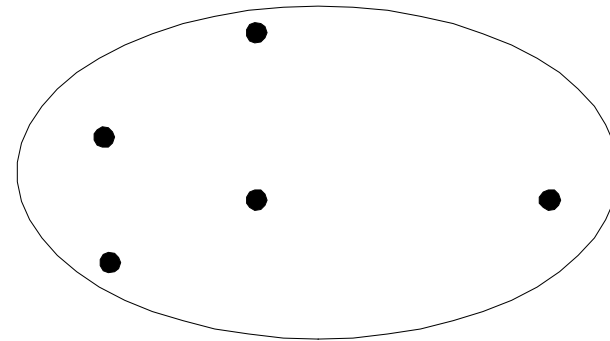
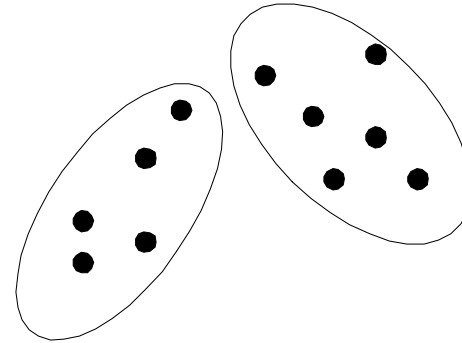


- Clustering
  - Grouping objects into a set of clusters based on similarity
- Clustering approaches
  - [Partitional clustering](#)
  - Hierarchical clustering
  - Conceptual clustering

# Partitional Clustering



**Original Points**

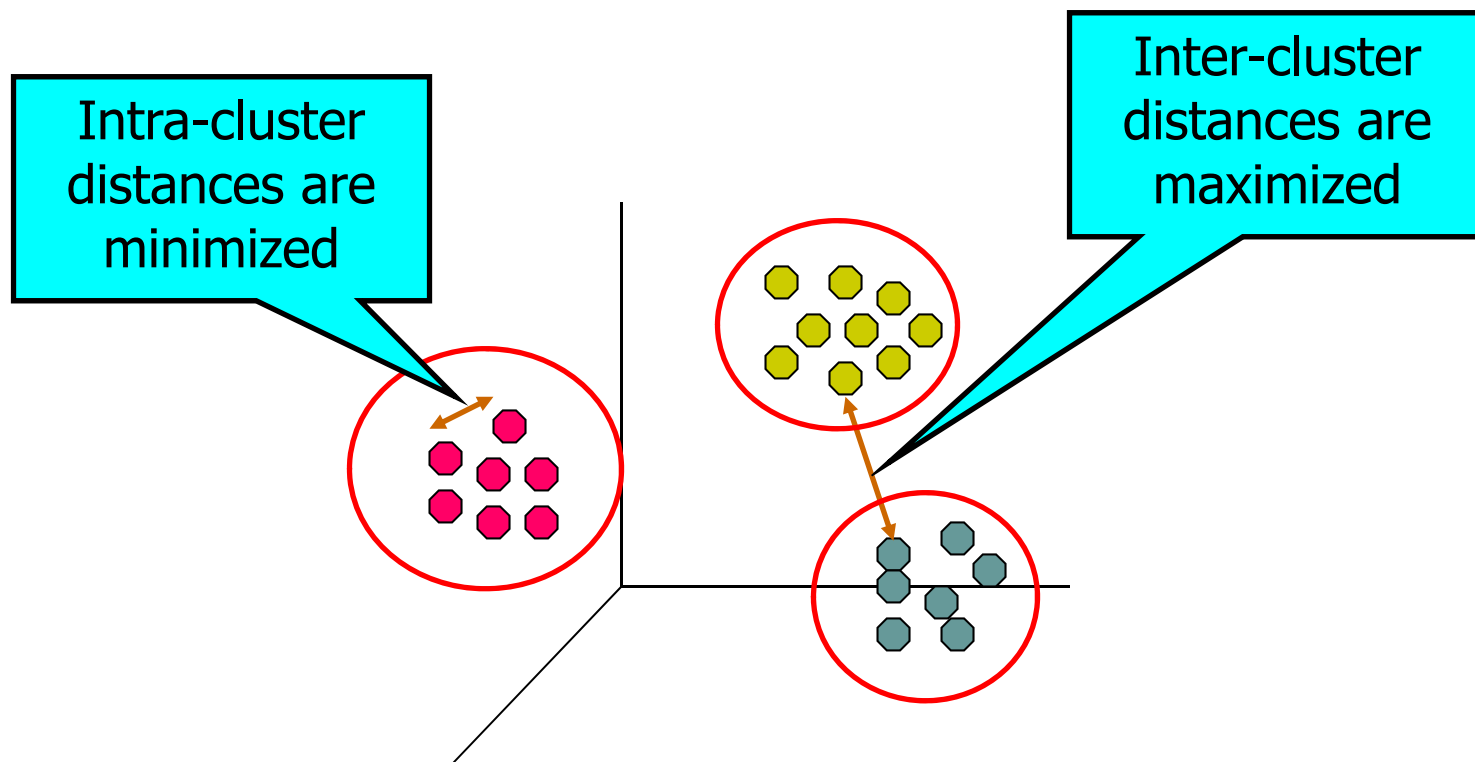


**A Partitional Clustering**



# Goal of Clustering

- Finding groups of objects such that the objects in a group will be similar to one another and different from the objects in other groups. It is a computationally difficult problem (NP-hard)





# Similarity and Dissimilarity



- **Similarity**

- Numerical measure of how alike two data objects are
- Is higher when objects are more alike
- Often falls in the range [0,1]

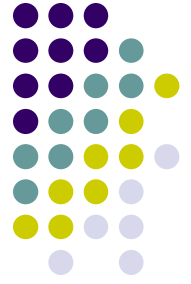
- **Dissimilarity**

- Numerical measure of how different two data objects are
- Lower when objects are more alike
- Minimum dissimilarity is often 0
- Upper limit varies

$$dist = \left( \sum_{k=1}^n |p_k - q_k|^r \right)^{\frac{1}{r}}$$

**Minkowski distance** is a generalization of **Euclidean distance**, where  $r$  is a parameter,  $n$  is the number of dimensions (attributes) and  $p_k$  and  $q_k$  are, respectively, the  $k$ th attributes (components) or data objects  $p$  and  $q$ .

# K-means Clustering



- **K-means algorithm:**
  - Partitional clustering approach
  - Each cluster is associated with a **centroid** (center point)
  - Each point is assigned to the cluster with the closest centroid
  - Number of clusters,  $K$ , must be specified as a parameter

- 
- 1: Select  $K$  points as the initial centroids.
  - 2: **repeat**
  - 3:   Form  $K$  clusters by assigning all points to the closest centroid.
  - 4:   Recompute the centroid of each cluster.
  - 5: **until** The centroids don't change
- 

+Several variations exist, e.g., fuzzy C-means, k-means++, etc.

# Unsupervised vs. Supervised Learning



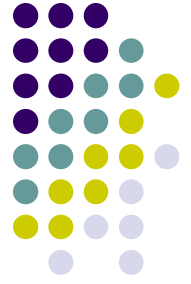
- **Applications for unsupervised learning**
  - Used in many disciplines such as Biology, Business, etc.
  - Pattern or object detection recognition: How can we group pixels into meaning regions of objects?
  - Outlier detection
  - Dimensionality reduction
  - Identification of common properties of objects, e.g., consumers
- **Some questions to clarify**
  - What's the results of unsupervised learning (clustering)?
  - What's the results of supervised learning?
  - How can we use of the results of unsupervised learning for supervised learning?
  - Is unsupervised learning more powerful than supervised learning (in terms of usefulness and accuracy of learned models)?
- **How can we combine both?**
  - Semi-supervised learning (like human learning)



# Domain Knowledge in Learning

- Use of domain knowledge in learning
  - Human can learn patterns by only a few examples quickly.
  - Cognitive scientist argue that learners learn better and faster based on the existing domain knowledge.
  - Most machine learning algorithms are designed to learn concepts without using domain knowledge.
- **Meta-DENDRAL** (by Buchanan and Mitchel, 1978) used domain knowledge
  - Knowledge is represented in the form of rules (based on graphs of molecular structure for LHS and RHS)
  - Meta-DENDRAL infers the rules for the structure of organic molecules from their chemical formula and data.
  - Explanations of the results are added after the inference as positive examples for a rule.

# Reinforcement Learning

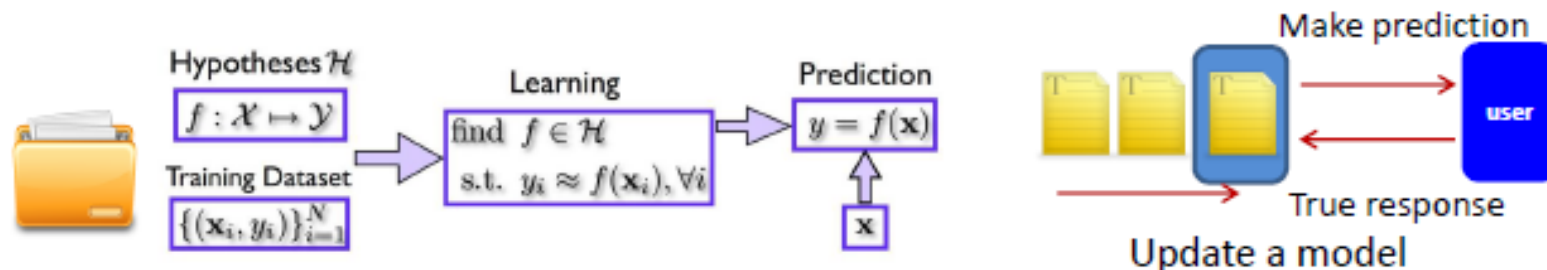


- Unlike supervised learning, **learning agent is not told directly what to do**, instead it goes through **trial-and-error**, **receives feedback or reward**, and learns an **optimal policy** to accomplish goals within its environment.
  - Inspired by old behaviorist psychology, concerned with how an agent ought to take actions in an environment so as to maximize some notion of cumulative reward.
  - Reinforcement learning is most successful when the environment is big or cannot be precisely described.
- Studied in many other disciplines
  - Control theory, operations research (called approximate dynamic programming), information theory, simulation-based optimization
  - In economics and game theory, reinforcement learning may be used to explain how equilibrium may arise under bounded rationality.



# Online Learning as a Promising Method for Big Data Analysis

- Batch/Offline learning
  - Observe a **batch** of training data  $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$
  - Learn a model from them
  - Predict new samples accurately
- Online learning
  - Observe a **sequence** of data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_t, y_t)$
  - Learn a model **incrementally** as instances come
  - Make the sequence of online predictions accurately





# Categories of Machine Learning

- Symbolic learning vs. non-symbolic learning
- **Supervised learning vs. unsupervised learning**
- **Batch learning vs. Online learning**
  - Difference between online learning and reinforcement learning is very small.
- **Other types of learning**
  - Analogical learning
  - Deep learning
  - Reinforcement learning
    - Important to robot
  - Hybrid learning or universal learning
    - Important for deep learning



# Applications of Machine Learning

- Data mining
  - Business intelligence, Big data analytics
- Financial modeling
  - Credit risk analysis, loan application screening
- Weather forecasting
- Speech recognition
- Games
- Many engineering problem solving
- Robotics
- Medicine and biology
- So many others

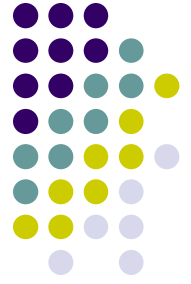




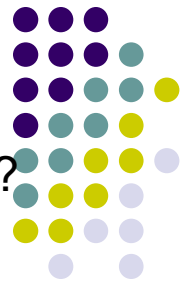
# Challenges of Machine Learning

- Learning from big data
  - Online learning
- Integration of various machine learning
  - Ensemble learning
  - Deep learning
  - Universal learner
- Inconsistency in knowledge representation
- Knowledge integration and utilization
  - Use of domain knowledge
  - Use of common sense knowledge

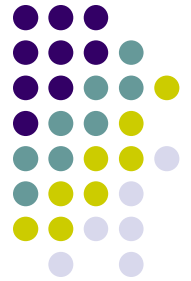
# Review Questions



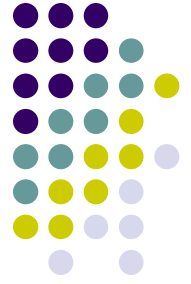
- What is the principle idea used in non-symbolic learning approaches?
- What are examples of non-symbolic learning algorithms?
- How does Genetic Learning method work?
- Why may the accuracy of most non-symbolic learning algorithms not be impacted by a few bad examples in a training data set?
- What is artificial neuron?
- What are the key components of artificial neuron? How does it work?
- How are input signals processed in an artificial neuron?
- How is the output of an artificial neuron determined?
- What is Artificial Neural Network (ANN)?
- How can the McCulloch-Pitts neuron model solve a problem (e.g., AND, OR)?
- What's the primary limitation of the McCulloch-Pitts neuron model?
- How does Perceptron Learning by Rosenblatt work?
- How can the Perceptron Learning algorithm learn the weights for a training data set? Try some examples.
- How can the learned weights be used for solving a problem? What are the applications?



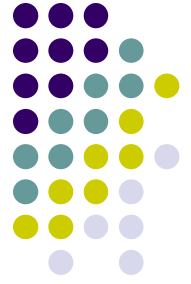
- What is the primary limitation of the Perceptron Learning by Rosenblatt?
- How can we overcome the limitation of the Perceptron Learning?
- What is sigmoidal function? What are the benefits?
- How is the output of a neuron using sigmoidal function determined?
- How does a multi-layer neural network with backpropagation work?
- What is the principle idea of the Delta rule by Rumelhart? Derive the Delta rule.
- How do we determine addition (+) or subtraction (-) of  $\Delta W$  in weight adjustment?
- Derive the formula of  $\Delta W$  for output nodes in a multi-layer neural net.
- Derive the formula of  $\Delta W$  for hidden nodes in a multi-layer neural net.
- Why is the formula for  $\Delta W$  for output nodes and hidden nodes different?
- Try some example calculation of  $\Delta W$  for both output nodes and hidden nodes.
- How can we implement a Deep Learning system based on the ANN approach?
- What is unsupervised learning? How can we implement an unsupervised learning?



- What's the primary difference between supervised learning and unsupervised learning?
- How does a clustering algorithm work in general? What are the necessary components for a clustering algorithm?
- What's the concept of "cluster"? What's the difference between "cluster" and "class"?
- Is it always possible to find the correct number of clusters in a data set?
- How can a similarity be measured between two objects?
- How does K-means work? Try some examples using K-means algorithm.
- What are the applications of clustering or unsupervised learning?
- What's the outcome of supervised learning?
- What's the outcome of unsupervised learning?
- Is supervised learning more useful than unsupervised learning?
- How can we use both supervised learning and unsupervised learning methods?
- What's the main idea of reinforcement learning? How is it different from other learning method such as supervised learning?
- Why is domain knowledge important in machine learning?



- What's the main idea of batch/offline learning?
- What's the main idea of online learning? Why is it nowadays important?
- What's the main idea of deep learning?
- What's the main idea of universal learning?
- Why is integrating various machine learning methods difficult?
- What are the applications of machine learning?
- What are the challenges of machine learning?



# References

- George Fluger, Artificial Intelligence: Structures and Strategies for Complex Problem Solving, 6<sup>th</sup> edition, **Chapters 10 and 12**, Addison Wesley, 2009.